



INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

Introducing an AR Model

Rob Reider

Adjunct Professor, NYU-Courant
Consultant, Quantopian

Mathematical Description of AR(1) Model

today value = mean + phi * yesterday's value + noise

$$R_t = \mu + \phi R_{t-1} + \epsilon_t$$

- Since only one lagged value on right hand side, this is called:
 - AR model of order 1, or
 - AR(1) model
- AR parameter is $\phi = 1$ (random walk), $= 0$ (white noise)
- For stationarity, $-1 < \phi < 1$ (to be stable and stationary)



Interpretation of AR(1) Parameter

R_t is a time series of stock returns


$$R_t = \mu + \phi R_{t-1} + \epsilon_t$$

- Negative ϕ : Mean Reversion
phi is negative, a positive return last period, at time t-1 implies that this period's return is more likely to be negative
- Positive ϕ : Momentum
a positive return last period implies that this period's return is expected to be positive

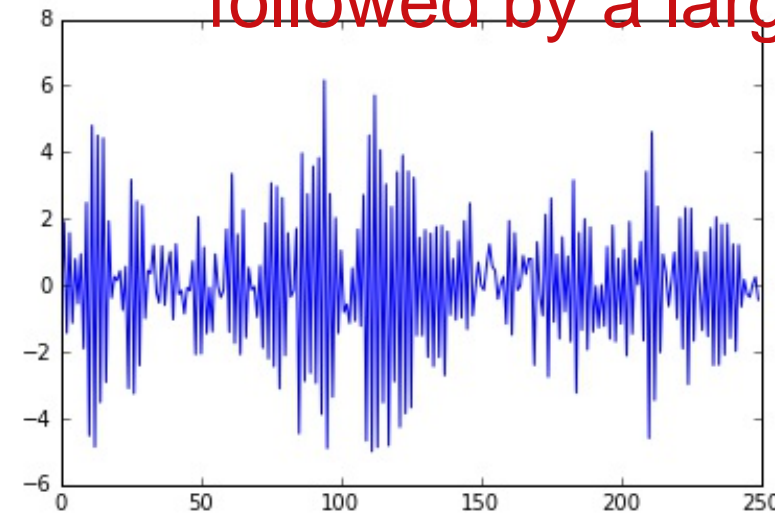


Comparison of AR(1) Time Series

- $\phi = 0.9$ close to random walk




- $\phi = -0.9$ more erratic (a large + value is usually followed by a large - value)




bottom two are similar, but are less exaggerated and closer to white noise

- $\phi = 0.5$



- $\phi = -0.5$





Comparison of AR(1) Autocorrelation Functions

autocorrelation decays exponentially


at a rate of phi

- $\phi = 0.9$

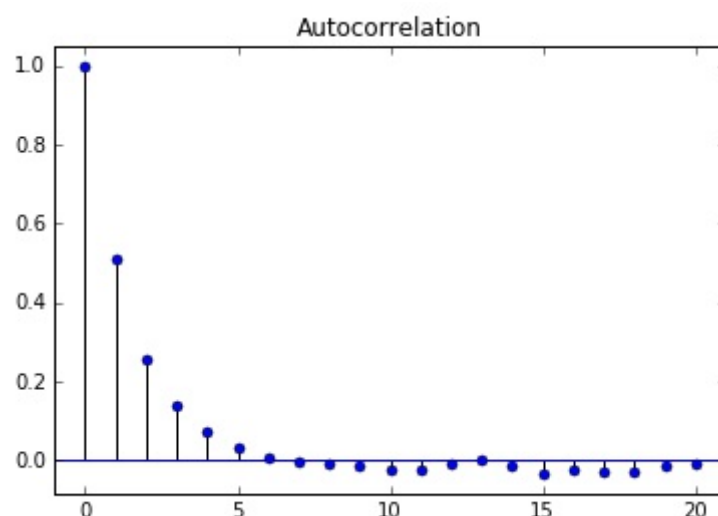
lag-1 auto
correlation is
0.9

lag-2 is 0.9^2

lag-3 is 0.9^3




- $\phi = 0.5$




autocorrelation function is still decays exponentially, but the signs of acf reverse at each lag

- $\phi = -0.9$



- $\phi = -0.5$





Higher Order AR Models

- AR(1)

can be extended to include more lagged values and more phi parameters.

$$R_t = \mu + \phi_1 R_{t-1} + \epsilon_t$$

- AR(2)

$$R_t = \mu + \phi_1 R_{t-1} + \phi_2 R_{t-2} + \epsilon_t$$

- AR(3)

$$R_t = \mu + \phi_1 R_{t-1} + \phi_2 R_{t-2} + \phi_3 R_{t-3} + \epsilon_t$$

- ...



Simulating an AR Process

want to study and understand a pure AR process,
it is useful to work with simulated data

```
from statsmodels.tsa.arima_process import ArmaProcess
ar = np.array([1, -0.9])
ma = np.array([1])
AR_object = ArmaProcess(ar, ma)
simulated_data = AR_object.generate_sample(nsample=1000)
plt.plot(simulated_data)
```

define order and parameters

the convention is a little counterintuitive: must include the zero-lag coeff of 1, and sign of the other coeff is the opposite of what we have been using.

EX: AR(1) with $\phi=0.9$ --> 2nd element of 'ar' array should be -0.9

this is consistent with the time series literature in the field of signal processing

also have to input MA parameters



INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

Let's practice!



INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

Estimating and Forecasting an AR Model

Rob Reider

Adjunct Professor, NYU-Courant
Consultant, Quantopian



Estimating an AR Model

- To estimate parameters from data (simulated)

```
from statsmodels.tsa.arima_model import ARMA
mod = ARMA(simulated_data, order=(1,0))
result = mod.fit()
```

create an instance of class called mod (data, order of the model)

(1,0) mean --> fit the data to AR(1)

(2,0) --> AR(2)

2nd part of order is MA part

.fit() to estimate model

Estimating an AR Model

- Full output (true $\mu = 0$ and $\phi = 0.9$)

estimated parameters are very close to true parameters

```
print(result.summary())
```

```
=====
                        ARMA Model Results
=====
Dep. Variable:          y      No. Observations:      5000
Model:                 ARMA(1, 0)  Log Likelihood      -7178.386
Method:                css-mle   S.D. of innovations      1.017
Date:                 Fri, 01 Dec 2017  AIC              14362.772
Time:                 15:34:50    BIC              14382.324
Sample:              0          HQIC              14369.625
=====
```

	coef	std err	z	P> z	[95.0% Conf. Int.]	
const	-0.0361	0.152	-0.238	0.812	-0.333	0.261
ar.L1.y	0.9054	0.006	151.020	0.000	0.894	0.917

```
=====
                        Roots
=====
```

	Real	Imaginary	Modulus	Frequency
AR.1	1.1045	+0.0000j	1.1045	0.0000

```
=====
```

mean mu

AR(1) parameter



Estimating an AR Model

- Only the estimates of μ and ϕ (true $\mu = 0$ and $\phi = 0.9$)

```
print(result.params)  see only coeff  
array([-0.03605989,  0.90535667])
```


μ

ϕ



Forecasting an AR Model

```
from statsmodels.tsa.arima_model import ARMA
mod = ARMA(simulated_data, order=(1,0))
res = mod.fit()
res.plot_predict(start='2016-07-01', end='2017-06-01') to do forecast
plt.show()
```



plot also gives confidence intervals around the out of sample forecasts.

notice how the confidence interval gets wider, the farther out the forecast is



INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

Let's practice!



INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

Choosing the Right Model

Rob Reider

Adjunct Professor, NYU-Courant
Consultant, Quantopian

in practice, you will ordinarily not be told the order of the model that you are trying to estimate



Identifying the Order of an AR Model

- The order of an $AR(p)$ model will usually be unknown
- Two techniques to determine order
 - Partial Autocorrelation Function
 - Information criteria

Partial Autocorrelation Function (PACF)

measures the incremental benefit of adding another lag.

$$R_t = \phi_{0,1} + \boxed{\phi_{1,1}} R_{t-1} + \epsilon_{1t}$$

imagine running several regressions, where u regress returns on more and more lagged values.

$$R_t = \phi_{0,2} + \phi_{1,2} R_{t-1} + \boxed{\phi_{2,2}} R_{t-2} + \epsilon_{2t}$$

$$R_t = \phi_{0,3} + \phi_{1,3} R_{t-1} + \phi_{2,3} R_{t-2} + \boxed{\phi_{3,3}} R_{t-3} + \epsilon_{3t}$$

$$R_t = \phi_{0,4} + \phi_{1,4} R_{t-1} + \phi_{2,4} R_{t-2} + \phi_{3,4} R_{t-3} + \boxed{\phi_{4,4}} R_{t-4} + \epsilon_{4t}$$

\vdots

is the lag-4 value of the partial acf, it represent how significant adding a fourth lag is when u already have 3 lags

coefficients in red boxes represent values of the partial autocorrelation function for different lags.

Plot PACF in Python

- Same as ACF, but use `plot_pacf` instead of `plt_acf`
- Import module

```
from statsmodels.graphics.tsaplots import plot_pacf
```

- Plot the PACF

```
plot_pacf(x, lags= 20, alpha=0.05)
```

how many lags of pacf will be plotted


alpha: set the width of the confidence interval



Comparison of PACF for Different AR Models


- AR(1)

only lag-1
pacf is
significantly
different
from zero




- AR(2)

2 lags are different
from zero




- AR(3)

3 lags are
different
from zero



- White Noise

no lags are significantly
different from zero





Information Criteria

more parameters in a model, better the model will fit the data.
but this can lead to overfitting of the data

- Information criteria: adjusts goodness-of-fit for number of parameters
- Two popular adjusted goodness-of-fit measures
 - AIC (Akaike Information Criterion)
 - BIC (Bayesian Information Criterion)

by imposing a penalty based on
the no of parameters used.



Information Criteria

- Estimation output

```
=====
                        ARMA Model Results
=====
Dep. Variable:          y      No. Observations:      2500
Model:                 ARMA(2, 0)  Log Likelihood      -3536.481
Method:                css-mle   S.D. of innovations    0.996
Date:                  Fri, 29 Dec 2017  AIC              7080.963
Time:                  22:53:24      BIC              7104.259
Sample:                0      HQIC              7089.420
=====

              coef      std err          z      P>|z|      [95.0% Conf. Int.]
-----
const          0.0054      0.010         0.517      0.605      -0.015      0.026
ar.L1.y        -0.6130      0.019       -32.243      0.000      -0.650     -0.576
ar.L2.y        -0.3109      0.019       -16.351      0.000      -0.348     -0.274
=====
                        Roots
=====
              Real      Imaginary      Modulus      Frequency
-----
AR.1          -0.9859      -1.4982j          1.7935      -0.3426
AR.2          -0.9859      +1.4982j          1.7935       0.3426
=====
```

Getting Information Criteria From statsmodels

- You learned earlier how to fit an AR model

```
from statsmodels.tsa.arima_model import ARMA
mod = ARMA(simulated_data, order=(1,0))
result = mod.fit()
```

- And to get full output

```
result.summary()
```

- Or just the parameters

```
result.params
```


- To get the AIC and BIC

```
result.aic
result.bic
```

Information Criteria

- Fit a simulated AR(3) to different AR(p) models
- Choose p with the lowest BIC

in practice, the way to use the info criteria is to fit several models, each with a different no of parameters, and choose the one with the lowest BIC





INTRODUCTION TO TIME SERIES ANALYSIS IN PYTHON

Let's practice!